

Colab Link:

[https://colab.research.google.com/drive/1_MAJ4RMmeow8ldscbCY47MxOODlnskES?
usp=sharing](https://colab.research.google.com/drive/1_MAJ4RMmeow8ldscbCY47MxOODlnskES?usp=sharing)

- **Installation of pandas**
 - Importing pandas
 - Importing the dataset
 - Dataframe/Series
- **Basic ops on a DataFrame**
 - df.info()
 - df.head()
 - df.tail()
 - df.shape()
 - df.describe()
- **Basic ops on columns**
 - Different ways of accessing cols
 - Check for Unique values
 - Rename column
 - Deleting col
 - Creating new cols
 - Quiz1 added
- **Basic ops on rows**
 - Implicit/explicit index
 - df.index[]
 - Indexing in series
 - Slicing in series
 - loc/iloc
 - Indexing/Slicing in dataframe
 - Adding a row
 - Check for duplicates
 - Deleting a row
- **Working with both rows and cols**
 - Quiz2 added
- **More in-built ops in pandas**
 - sum()
 - count()
 - mean()
- **Sorting**
 - Quiz3 added
- **Creating series and Dataframes from scratch**

```
!pip install pandas
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-w>

Requirement already satisfied: pandas in /usr/local/lib/python3.7/dist-package
 Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.7/dist-package
 Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.7/dist-package
 Requirement already satisfied: pytz>=2017.3 in /usr/local/lib/python3.7/dist-package
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-package

```
import pandas as pd
import numpy as np
```


```
!wget "https://drive.google.com/uc?export=download&id=1E3bwvYGf1ig32RmcYiWc0I1"
```

```
--2022-06-28 15:41:33-- https://drive.google.com/uc?export=download&id=1E3bwvYGf1ig32RmcYiWc0I1
Resolving drive.google.com (drive.google.com)... 142.251.107.102, 142.251.107.102
Connecting to drive.google.com (drive.google.com)|142.251.107.102|:443... conn
HTTP request sent, awaiting response... 303 See Other
Location: https://doc-0s-68-docs.googleusercontent.com/docs/securesc/ha0ro937c
Warning: wildcards not supported in HTTP.
--2022-06-28 15:41:34-- https://doc-0s-68-docs.googleusercontent.com/docs/securesc/ha0ro937c
Resolving doc-0s-68-docs.googleusercontent.com (doc-0s-68-docs.googleusercontent.com)... 142.251.107.102
Connecting to doc-0s-68-docs.googleusercontent.com (doc-0s-68-docs.googleusercontent.com)|142.251.107.102|:443... conn
HTTP request sent, awaiting response... 200 OK
Length: 83785 (82K) [text/csv]
Saving to: 'mckinsey.csv'
```

```
mckinsey.csv          100%[=====>]  81.82K  --.-KB/s    in 0.001s
```

```
2022-06-28 15:41:34 (88.4 MB/s) - 'mckinsey.csv' saved [83785/83785]
```

```
df = pd.read_csv("mckinsey.csv")
df
```

	country	year	population	continent	life_exp	gdp_cap	
0	Afghanistan	1952	8425333	Asia	28.801	779.445314	
1	Afghanistan	1957	9240934	Asia	30.332	820.853030	
2	Afghanistan	1962	10267083	Asia	31.997	853.100710	
3	Afghanistan	1967	11537966	Asia	34.020	836.197138	
4	Afghanistan	1972	13079460	Asia	36.088	739.981106	
...	
1699	Zimbabwe	1987	9216418	Africa	62.351	706.157306	
1700	Zimbabwe	1992	10704340	Africa	60.377	693.420786	
1701	Zimbabwe	1997	11404948	Africa	46.809	792.449960	
1702	Zimbabwe	2002	11926563	Africa	39.989	672.038623	
1703	Zimbabwe	2007	12311143	Africa	43.487	469.709298	

1704 rows x 6 columns

```
type(df)
```

pandas.core.frame.DataFrame

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1704 entries, 0 to 1703
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   country         1704 non-null   object
1   year            1704 non-null   int64
2   population       1704 non-null   int64
3   continent       1704 non-null   object
4   life_exp        1704 non-null   float64
5   gdp_cap         1704 non-null   float64
dtypes: float64(2), int64(2), object(2)
memory usage: 80.0+ KB
```


```
df.shape
```

(1704, 6)


```
df.head(15)
```

	country	year	population	continent	life_exp	gdp_cap	
0	Afghanistan	1952	8425333	Asia	28.801	779.445314	
1	Afghanistan	1957	9240934	Asia	30.332	820.853030	
2	Afghanistan	1962	10267083	Asia	31.997	853.100710	
3	Afghanistan	1967	11537966	Asia	34.020	836.197138	
4	Afghanistan	1972	13079460	Asia	36.088	739.981106	
5	Afghanistan	1977	14880372	Asia	38.438	786.113360	
6	Afghanistan	1982	12881816	Asia	39.854	978.011439	
7	Afghanistan	1987	13867957	Asia	40.822	852.395945	
8	Afghanistan	1992	16317921	Asia	41.674	649.341395	
9	Afghanistan	1997	22227415	Asia	41.763	635.341351	
10	Afghanistan	2002	25268405	Asia	42.129	726.734055	
11	Afghanistan	2007	31889923	Asia	43.828	974.580338	
12	Albania	1952	1282697	Europe	55.230	1601.056136	
13	Albania	1957	1476505	Europe	59.280	1942.284244	
14	Albania	1962	1728137	Europe	64.820	2312.888958	


```
df.tail()
```

	country	year	population	continent	life_exp	gdp_cap	
1699	Zimbabwe	1987	9216418	Africa	62.351	706.157306	
1700	Zimbabwe	1992	10704340	Africa	60.377	693.420786	
1701	Zimbabwe	1997	11404948	Africa	46.809	792.449960	
1702	Zimbabwe	2002	11926563	Africa	39.989	672.038623	
1703	Zimbabwe	2007	12311143	Africa	43.487	469.709298	

```
df.describe()
```

	year	population	life_exp	gdp_cap	
count	1704.00000	1.704000e+03	1704.000000	1704.000000	
mean	1979.50000	2.960121e+07	59.474439	7215.327081	
std	17.26533	1.061579e+08	12.917107	9857.454543	
min	1952.00000	6.001100e+04	23.599000	241.165876	
25%	1965.75000	2.793664e+06	48.198000	1202.060309	
50%	1979.50000	7.023596e+06	60.712500	3531.846988	
75%	1993.25000	1.958522e+07	70.845500	9325.462346	
max	2007.00000	1.318683e+09	82.603000	113523.132900	

```
df.describe(include="object")
```

	country	continent	
count	1704	1704	
unique	142	5	
top	Afghanistan	Africa	
freq	12	624	

```
# Basic operations on columns
```

```
df.columns
```

```
Index(['country', 'year', 'population', 'continent', 'life_exp', 'gdp_cap'],  
      dtype='object')
```

```
df.keys()
```

```
Index(['country', 'year', 'population', 'continent', 'life_exp', 'gdp_cap'],
      dtype='object')
```

```
df["country"]
```

```
0      Afghanistan
1      Afghanistan
2      Afghanistan
3      Afghanistan
4      Afghanistan
...
1699    Zimbabwe
1700    Zimbabwe
1701    Zimbabwe
1702    Zimbabwe
1703    Zimbabwe
Name: country, Length: 1704, dtype: object
```

```
type(df["country"])
```

```
pandas.core.series.Series
```

```
df[["country", "continent"]].head()
```

	country	continent
0	Afghanistan	Asia
1	Afghanistan	Asia
2	Afghanistan	Asia
3	Afghanistan	Asia
4	Afghanistan	Asia

```
df["country"].unique()
```

```
array(['Afghanistan', 'Albania', 'Algeria', 'Angola', 'Argentina',
      'Australia', 'Austria', 'Bahrain', 'Bangladesh', 'Belgium',
      'Benin', 'Bolivia', 'Bosnia and Herzegovina', 'Botswana', 'Brazil',
      'Bulgaria', 'Burkina Faso', 'Burundi', 'Cambodia', 'Cameroon',
      'Canada', 'Central African Republic', 'Chad', 'Chile', 'China',
      'Colombia', 'Comoros', 'Congo, Dem. Rep.', 'Congo, Rep.',
      'Costa Rica', 'Cote d'Ivoire', 'Croatia', 'Cuba', 'Czech Republic',
      'Denmark', 'Djibouti', 'Dominican Republic', 'Ecuador', 'Egypt',
      'El Salvador', 'Equatorial Guinea', 'Eritrea', 'Ethiopia',
      'Finland', 'France', 'Gabon', 'Gambia', 'Germany', 'Ghana',
      'Greece', 'Guatemala', 'Guinea', 'Guinea-Bissau', 'Haiti',
      'Honduras', 'Hong Kong, China', 'Hungary', 'Iceland', 'India',
      'Indonesia', 'Iran', 'Iraq', 'Ireland', 'Israel', 'Italy',
      'Jamaica', 'Japan', 'Jordan', 'Kenya', 'Korea, Dem. Rep.',
      'Korea, Rep.', 'Kuwait', 'Lebanon', 'Lesotho', 'Liberia', 'Libya',
      'Madagascar', 'Malawi', 'Malaysia', 'Mali', 'Mauritania',
      'Mauritius', 'Mexico', 'Mongolia', 'Montenegro', 'Morocco',
      'Mozambique', 'Myanmar', 'Namibia', 'Nepal', 'Netherlands',
```

```
'New Zealand', 'Nicaragua', 'Niger', 'Nigeria', 'Norway', 'Oman',
'Pakistan', 'Panama', 'Paraguay', 'Peru', 'Philippines', 'Poland',
'Portugal', 'Puerto Rico', 'Reunion', 'Romania', 'Rwanda',
'Sao Tome and Principe', 'Saudi Arabia', 'Senegal', 'Serbia',
'Sierra Leone', 'Singapore', 'Slovak Republic', 'Slovenia',
'Somalia', 'South Africa', 'Spain', 'Sri Lanka', 'Sudan',
'Swaziland', 'Sweden', 'Switzerland', 'Syria', 'Taiwan',
'Tanzania', 'Thailand', 'Togo', 'Trinidad and Tobago', 'Tunisia',
'Turkey', 'Uganda', 'United Kingdom', 'United States', 'Uruguay',
'Venezuela', 'Vietnam', 'West Bank and Gaza', 'Yemen, Rep.',
'Zambia', 'Zimbabwe'], dtype=object)
```

```
len(df["country"].unique())
```

```
142
```

```
df['country'].value_counts()
```

```
Afghanistan      12
Pakistan          12
New Zealand       12
Nicaragua         12
Niger             12
..
Eritrea           12
Equatorial Guinea 12
El Salvador       12
Egypt             12
Zimbabwe          12
Name: country, Length: 142, dtype: int64
```

```
df["country"].nunique()
```

```
142
```

```
df.rename({"country": "Country"}, axis=1, inplace=True)
```

```
df
```

	Country	year	population	continent	life_exp	gdp_cap
0	Afghanistan	1952	8425333	Asia	28.801	779.445314
1	Afghanistan	1957	9240934	Asia	30.332	820.853030
2	Afghanistan	1962	10267083	Asia	31.997	853.100710
3	Afghanistan	1967	11537966	Asia	34.020	836.197138



```
df.rename({"Country":"country"}, axis=1, inplace=True)
```

```
df.country # dont use it
```

0	Afghanistan
1	Afghanistan
2	Afghanistan
3	Afghanistan
4	Afghanistan
...	
1699	Zimbabwe
1700	Zimbabwe
1701	Zimbabwe
1702	Zimbabwe
1703	Zimbabwe

Name: country, Length: 1704, dtype: object

```
df.drop("continent", axis=1, inplace=True)
```

```
df
```


	country	year	population	life_exp	gdp_cap
0	Afghanistan	1952	8425333	28.801	779.445314
1	Afghanistan	1957	9240934	30.332	820.853030
2	Afghanistan	1962	10267083	31.997	853.100710
3	Afghanistan	1967	11537966	34.020	836.197138
4	Afghanistan	1972	13079460	36.088	739.981106
...
1699	Zimbabwe	1987	9216418	62.351	706.157306
1700	Zimbabwe	1992	10704340	60.377	693.420786
1701	Zimbabwe	1997	11404948	46.809	792.449960
1702	Zimbabwe	2002	11926563	39.989	672.038623
1703	Zimbabwe	2007	12311143	43.487	469.709298



1704 rows x 5 columns

```
df["New"] = df["life_exp"] + df["year"]
```



df

	country	year	population	life_exp	gdp_cap	New	
0	Afghanistan	1952	8425333	28.801	779.445314	1980.801	
1	Afghanistan	1957	9240934	30.332	820.853030	1987.332	
2	Afghanistan	1962	10267083	31.997	853.100710	1993.997	
3	Afghanistan	1967	11537966	34.020	836.197138	2001.020	
4	Afghanistan	1972	13079460	36.088	739.981106	2008.088	
...	
1699	Zimbabwe	1987	9216418	62.351	706.157306	2049.351	
1700	Zimbabwe	1992	10704340	60.377	693.420786	2052.377	
1701	Zimbabwe	1997	11404948	46.809	792.449960	2043.809	
1702	Zimbabwe	2002	11926563	39.989	672.038623	2041.989	
1703	Zimbabwe	2007	12311143	43.487	469.709298	2050.487	

1704 rows x 6 columns

```
df['Own'] = [i for i in range(df.shape[0])]
```

df

	country	year	population	life_exp	gdp_cap	New	Own	
0	Afghanistan	1952	8425333	28.801	779.445314	1980.801	0	
1	Afghanistan	1957	9240934	30.332	820.853030	1987.332	1	
2	Afghanistan	1962	10267083	31.997	853.100710	1993.997	2	
3	Afghanistan	1967	11537966	34.020	836.197138	2001.020	3	
4	Afghanistan	1972	13079460	36.088	739.981106	2008.088	4	
...	
1699	Zimbabwe	1987	9216418	62.351	706.157306	2049.351	1699	
1700	Zimbabwe	1992	10704340	60.377	693.420786	2052.377	1700	
1701	Zimbabwe	1997	11404948	46.809	792.449960	2043.809	1701	
1702	Zimbabwe	2002	11926563	39.989	672.038623	2041.989	1702	
1703	Zimbabwe	2007	12311143	43.487	469.709298	2050.487	1703	

1704 rows x 7 columns

```
df.drop(["New", "Own"], axis=1, inplace=True)
```

df

	country	year	population	life_exp	gdp_cap
0	Afghanistan	1952	8425333	28.801	779.445314
1	Afghanistan	1957	9240934	30.332	820.853030
2	Afghanistan	1962	10267083	31.997	853.100710
3	Afghanistan	1967	11537966	34.020	836.197138
4	Afghanistan	1972	13079460	36.088	739.981106
...
1699	Zimbabwe	1987	9216418	62.351	706.157306
1700	Zimbabwe	1992	10704340	60.377	693.420786
1701	Zimbabwe	1997	11404948	46.809	792.449960
1702	Zimbabwe	2002	11926563	39.989	672.038623
1703	Zimbabwe	2007	12311143	43.487	469.709298

1704 rows x 5 columns

Working with Rows

ser = df["country"]

ser[4]

'Afghanistan'

ser[6:15]

```

6    Afghanistan
7    Afghanistan
8    Afghanistan
9    Afghanistan
10   Afghanistan
11   Afghanistan
12      Albania
13      Albania
14      Albania

```

Name: country, dtype: object

ser

```

0    Afghanistan
1    Afghanistan
2    Afghanistan
3    Afghanistan
4    Afghanistan

```

```

...
1699    Zimbabwe
1700    Zimbabwe
1701    Zimbabwe
1702    Zimbabwe
1703    Zimbabwe
Name: country, Length: 1704, dtype: object

```

```
ser.index = np.arange(1, df.shape[0]+1, step=1)
```

```
ser
```

```

1      Afghanistan
2      Afghanistan
3      Afghanistan
4      Afghanistan
5      Afghanistan
...
1700    Zimbabwe
1701    Zimbabwe
1702    Zimbabwe
1703    Zimbabwe
1704    Zimbabwe
Name: country, Length: 1704, dtype: object

```

```
# Indexes/Labels
```

```
# Explicit Indexes - visble outside, can be anything - numbers, 0-N-1, 1-N, st
```

```
# Implicit Indexes - 0-N-1
```

```
ser
```

```

1      Afghanistan
2      Afghanistan
3      Afghanistan
4      Afghanistan
5      Afghanistan
...
1700    Zimbabwe
1701    Zimbabwe
1702    Zimbabwe
1703    Zimbabwe
1704    Zimbabwe
Name: country, Length: 1704, dtype: object

```

```
ser.index[2]
```

```
3
```

```
data = pd.Series(["a", "b", "c"], index=[1, 5, 3])
```

```
data
```

```

1      a
5      b

```

```
3      c
dtype: object
```

```
data[1] # indexing uses explicit indices
```

```
'a'
```

```
data[1:3] # slicing used implicit indices
```

```
5      b
3      c
dtype: object
```

```
data = pd.Series(["a", "b", "c"], index=["x", "y", "z"])
data
```

```
x      a
y      b
z      c
dtype: object
```

```
data = pd.Series(["a", "b", "c"], index=[1, 2, 2])
data
```

```
1      a
2      b
2      c
dtype: object
```

```
data[2]
```

```
2      b
2      c
dtype: object
```

```
# indexing uses explicit indices
```

```
# slicing uses implicit indices
```

```
# Indexers - loc and iloc
```

```
# loc - Allows indexing and slicing that always references the explicit index
```

```
data = pd.Series(["a", "b", "c"], index=[1, 2, 2])
data
```

```
1      a
2      b
2      c
dtype: object
```

```
data.loc[2]
```

```
2    b
2    c
dtype: object
```

```
data.loc[1]
```

```
'a'
```

```
data.loc[1:3]
```

```
1    a
2    b
2    c
dtype: object
```

```
data.loc[1:2]
```

```
1    a
2    b
2    c
dtype: object
```

```
# iloc - Allows indexing and slicing that always references the implicit index
```

```
data.iloc[1]
```

```
'b'
```

```
data.iloc[0]
```

```
'a'
```

```
data.loc[0] # indexing it gives an error
```

```

-----
KeyError                                Traceback (most recent call last)
/usr/local/lib/python3.7/dist-packages/pandas/core/indexes/base.py in
get_loc(self, key, method, tolerance)
    3360                 try:
-> 3361                     return self._engine.get_loc(casted_key)
    3362             except KeyError as err:

```

↕ 8 frames

```
KeyError: 0
```

```
data.iloc[0:2] # end is not inclusive just like in Python normal indexing
```

```

1    a
2    b
dtype: object

```

```
3361                 return self._engine.get_loc(casted_key)
```

```
data.loc[100:105] # slicing it doesnt it give error
```

```
Series([], dtype: object)
```

```
df
```

	country	year	population	life_exp	gdp_cap	
0	Afghanistan	1952	8425333	28.801	779.445314	
1	Afghanistan	1957	9240934	30.332	820.853030	
2	Afghanistan	1962	10267083	31.997	853.100710	
3	Afghanistan	1967	11537966	34.020	836.197138	
4	Afghanistan	1972	13079460	36.088	739.981106	
...	
1699	Zimbabwe	1987	9216418	62.351	706.157306	
1700	Zimbabwe	1992	10704340	60.377	693.420786	
1701	Zimbabwe	1997	11404948	46.809	792.449960	
1702	Zimbabwe	2002	11926563	39.989	672.038623	
1703	Zimbabwe	2007	12311143	43.487	469.709298	

1704 rows x 5 columns

```
df.loc[3]
```

```

country    Afghanistan
year        1967
population    11537966
life_exp      34.02
gdp_cap      836.197138
Name: 3, dtype: object

```

```
df.iloc[3]
```

```
country    Afghanistan
year       1967
population  11537966
life_exp   34.02
gdp_cap    836.197138
Name: 3, dtype: object
```

```
df.index = df.index + 1
```

```
df
```

	country	year	population	life_exp	gdp_cap
1	Afghanistan	1952	8425333	28.801	779.445314
2	Afghanistan	1957	9240934	30.332	820.853030
3	Afghanistan	1962	10267083	31.997	853.100710
4	Afghanistan	1967	11537966	34.020	836.197138
5	Afghanistan	1972	13079460	36.088	739.981106
...
1700	Zimbabwe	1987	9216418	62.351	706.157306
1701	Zimbabwe	1992	10704340	60.377	693.420786
1702	Zimbabwe	1997	11404948	46.809	792.449960
1703	Zimbabwe	2002	11926563	39.989	672.038623
1704	Zimbabwe	2007	12311143	43.487	469.709298

1704 rows x 5 columns

```
df.iloc[[1, 10, 100]]
```

	country	year	population	life_exp	gdp_cap
2	Afghanistan	1957	9240934	30.332	820.853030
11	Afghanistan	2002	25268405	42.129	726.734055
101	Bangladesh	1972	70759295	45.252	630.233627

```
df.loc[[1, 10, 100]]
```

country	year	population	life_exp	gdp_cap
---------	------	------------	----------	---------



```
df.iloc[-1]
```

```
country      Zimbabwe
year          2007
population    12311143
life_exp      43.487
gdp_cap       469.709298
Name: 1704, dtype: object
```

```
df.loc[-1]
```

```
-----
ValueError                                Traceback (most recent call last)
/usr/local/lib/python3.7/dist-packages/pandas/core/indexes/range.py in
get_loc(self, key, method, tolerance)
    384         try:
--> 385             return self._range.index(new_key)
    386         except ValueError as err:
```

```
ValueError: -1 is not in range
```

The above exception was the direct cause of the following exception:

```
-----
KeyError                                Traceback (most recent call last)
5 frames
/usr/local/lib/python3.7/dist-packages/pandas/core/indexes/range.py in
get_loc(self, key, method, tolerance)
    385         return self._range.index(new_key)
    386         except ValueError as err:
--> 387             raise KeyError(key) from err
    388         raise KeyError(key)
    389         return super().get_loc(key, method=method,
tolerance=tolerance)
```

```
KeyError: -1
```

```
df.index = df["country"]
```

```
df
```


	country	year	population	life_exp	gdp_cap	
	country					
	Afghanistan	Afghanistan	1952	8425333	28.801	779.445314
	Afghanistan	Afghanistan	1957	9240934	30.332	820.853030
	Afghanistan	Afghanistan	1962	10267083	31.997	853.100710
	Afghanistan	Afghanistan	1967	11537966	34.020	836.197138
	Afghanistan	Afghanistan	1972	13079460	36.088	739.981106

```
df = pd.read_csv("mckinsey.csv")
df
```

	country	year	population	continent	life_exp	gdp_cap	
0	Afghanistan	1952	8425333	Asia	28.801	779.445314	
1	Afghanistan	1957	9240934	Asia	30.332	820.853030	
2	Afghanistan	1962	10267083	Asia	31.997	853.100710	
3	Afghanistan	1967	11537966	Asia	34.020	836.197138	
4	Afghanistan	1972	13079460	Asia	36.088	739.981106	
...	
1699	Zimbabwe	1987	9216418	Africa	62.351	706.157306	
1700	Zimbabwe	1992	10704340	Africa	60.377	693.420786	
1701	Zimbabwe	1997	11404948	Africa	46.809	792.449960	
1702	Zimbabwe	2002	11926563	Africa	39.989	672.038623	
1703	Zimbabwe	2007	12311143	Africa	43.487	469.709298	

1704 rows x 6 columns


```
df_temp = df.set_index("country")
df_temp
```

	year	population	continent	life_exp	gdp_cap	
country						
Afghanistan	1952	8425333	Asia	28.801	779.445314	
Afghanistan	1957	9240934	Asia	30.332	820.853030	
Afghanistan	1962	10267083	Asia	31.997	853.100710	
Afghanistan	1967	11537966	Asia	34.020	836.197138	


```
df_temp.loc["Afghanistan"]
```

	year	population	continent	life_exp	gdp_cap	
country						
Afghanistan	1952	8425333	Asia	28.801	779.445314	
Afghanistan	1957	9240934	Asia	30.332	820.853030	
Afghanistan	1962	10267083	Asia	31.997	853.100710	
Afghanistan	1967	11537966	Asia	34.020	836.197138	
Afghanistan	1972	13079460	Asia	36.088	739.981106	
Afghanistan	1977	14880372	Asia	38.438	786.113360	
Afghanistan	1982	12881816	Asia	39.854	978.011439	
Afghanistan	1987	13867957	Asia	40.822	852.395945	
Afghanistan	1992	16317921	Asia	41.674	649.341395	
Afghanistan	1997	22227415	Asia	41.763	635.341351	
Afghanistan	2002	25268405	Asia	42.129	726.734055	
Afghanistan	2007	31889923	Asia	43.828	974.580338	

```
df
```

	country	year	population	continent	life_exp	gdp_cap	
0	Afghanistan	1952	8425333	Asia	28.801	779.445314	
1	Afghanistan	1957	9240934	Asia	30.332	820.853030	
2	Afghanistan	1962	10267083	Asia	31.997	853.100710	

```
data_dict = {'country': 'India', 'year': 2000, 'life_exp': 37.08, 'population': 13500000}
df.append(data_dict, ignore_index=True)
```

	country	year	population	continent	life_exp	gdp_cap	
0	Afghanistan	1952	8425333	Asia	28.801	779.445314	
1	Afghanistan	1957	9240934	Asia	30.332	820.853030	
2	Afghanistan	1962	10267083	Asia	31.997	853.100710	
3	Afghanistan	1967	11537966	Asia	34.020	836.197138	
4	Afghanistan	1972	13079460	Asia	36.088	739.981106	
...	
1700	Zimbabwe	1992	10704340	Africa	60.377	693.420786	
1701	Zimbabwe	1997	11404948	Africa	46.809	792.449960	
1702	Zimbabwe	2002	11926563	Africa	39.989	672.038623	
1703	Zimbabwe	2007	12311143	Africa	43.487	469.709298	
1704	India	2000	13500000	NaN	37.080	900.230000	

1705 rows x 6 columns

```
df.loc[len(df.index)] = ["India", 2000, 1350000, "Asia", 37.080, 900.230000]
```

df

	country	year	population	continent	life_exp	gdp_cap
0	Afghanistan	1952	8425333	Asia	28.801	779.445314
1	Afghanistan	1957	9240934	Asia	30.332	820.853030
2	Afghanistan	1962	10267083	Asia	31.997	853.100710

```
df.loc[len(df.index)] = {'country': 'India', 'year': 2000, 'life_exp':37.08, 'population':13500000, 'gdp_cap':900.23}
df
```

	country	year	population	continent	life_exp	gdp_cap
0	Afghanistan	1952	8425333	Asia	28.801	779.445314
1	Afghanistan	1957	9240934	Asia	30.332	820.853030
2	Afghanistan	1962	10267083	Asia	31.997	853.100710
3	Afghanistan	1967	11537966	Asia	34.020	836.197138
4	Afghanistan	1972	13079460	Asia	36.088	739.981106
...
1701	Zimbabwe	1997	11404948	Africa	46.809	792.449960
1702	Zimbabwe	2002	11926563	Africa	39.989	672.038623
1703	Zimbabwe	2007	12311143	Africa	43.487	469.709298
1704	India	2000	1350000	Asia	37.080	900.230000
1706	India	2000	13500000	NaN	37.080	900.230000

1706 rows x 6 columns

```
df.iloc[len(df.index)] = {'country': 'India', 'year': 2000, 'life_exp':37.08, 'population':13500000, 'gdp_cap':900.23}
```

```
-----
IndexError                                Traceback (most recent call last)
<ipython-input-115-6ba8a45c0f22> in <module>()
----> 1 df.iloc[len(df.index)] = {'country': 'India', 'year':
2000, 'life_exp':37.08, 'population':13500000, 'gdp_cap':900.23}

1 frames
/usr/local/lib/python3.7/dist-packages/pandas/core/indexing.py in
_has_valid_setitem_indexer(self, indexer)
    1459         elif is_integer(i):
    1460             if i >= len(ax):
-> 1461                 raise IndexError("iloc cannot enlarge its target
object")
    1462         elif isinstance(i, dict):
    1463             raise IndexError("iloc cannot enlarge its target
object")

IndexError: iloc cannot enlarge its target object
```

```
df.loc[1705] = {'country': 'India', 'year': 2000, 'life_exp':37.08, 'population
```

df

	country	year	population	continent	life_exp	gdp_cap
0	Afghanistan	1952	8425333	Asia	28.801	779.445314
1	Afghanistan	1957	9240934	Asia	30.332	820.853030
2	Afghanistan	1962	10267083	Asia	31.997	853.100710
3	Afghanistan	1967	11537966	Asia	34.020	836.197138
4	Afghanistan	1972	13079460	Asia	36.088	739.981106
...
1702	Zimbabwe	2002	11926563	Africa	39.989	672.038623
1703	Zimbabwe	2007	12311143	Africa	43.487	469.709298
1704	India	2000	1350000	Asia	37.080	900.230000
1706	India	2000	13500000	NaN	37.080	900.230000
1705	India	2000	13500000	NaN	37.080	900.230000

1707 rows x 6 columns

```
df.drop(1704, axis=0)
```

	country	year	population	continent	life_exp	gdp_cap
0	Afghanistan	1952	8425333	Asia	28.801	779.445314
1	Afghanistan	1957	9240934	Asia	30.332	820.853030
2	Afghanistan	1962	10267083	Asia	31.997	853.100710
3	Afghanistan	1967	11537966	Asia	34.020	836.197138
4	Afghanistan	1972	13079460	Asia	36.088	739.981106
...
1701	Zimbabwe	1997	11404948	Africa	46.809	792.449960
1702	Zimbabwe	2002	11926563	Africa	39.989	672.038623
1703	Zimbabwe	2007	12311143	Africa	43.487	469.709298
1706	India	2000	13500000	NaN	37.080	900.230000
1705	India	2000	13500000	NaN	37.080	900.230000

1706 rows x 6 columns

```
df.loc[df.duplicated() ]
```

```
Int64Index([1705], dtype='int64')
```

```
df.drop_duplicates(keep=False, inplace=True)
```

```
df.drop(1704, axis=0, inplace=True)
```

```
df
```

	country	year	population	continent	life_exp	gdp_cap
0	Afghanistan	1952	8425333	Asia	28.801	779.445314
1	Afghanistan	1957	9240934	Asia	30.332	820.853030
2	Afghanistan	1962	10267083	Asia	31.997	853.100710
3	Afghanistan	1967	11537966	Asia	34.020	836.197138
4	Afghanistan	1972	13079460	Asia	36.088	739.981106
...
1699	Zimbabwe	1987	9216418	Africa	62.351	706.157306
1700	Zimbabwe	1992	10704340	Africa	60.377	693.420786
1701	Zimbabwe	1997	11404948	Africa	46.809	792.449960
1702	Zimbabwe	2002	11926563	Africa	39.989	672.038623
1703	Zimbabwe	2007	12311143	Africa	43.487	469.709298


1704 rows x 6 columns

```
# Working with rows and columns
```


```
df.iloc[1:5, 1:4]
```

	year	population	continent
1	1957	9240934	Asia
2	1962	10267083	Asia
3	1967	11537966	Asia
4	1972	13079460	Asia

```
df.loc[1:5, "country":"continent"]
```

	country	year	population	continent	
1	Afghanistan	1957	9240934	Asia	
2	Afghanistan	1962	10267083	Asia	

```
df.iloc[1:10:2]
```

	country	year	population	continent	life_exp	gdp_cap	
1	Afghanistan	1957	9240934	Asia	30.332	820.853030	
3	Afghanistan	1967	11537966	Asia	34.020	836.197138	
5	Afghanistan	1977	14880372	Asia	38.438	786.113360	
7	Afghanistan	1987	13867957	Asia	40.822	852.395945	
9	Afghanistan	1997	22227415	Asia	41.763	635.341351	

```
# Built-in operations in pandas
```

```
df["life_exp"].mean()
```

```
59.474439366197174
```

```
df["life_exp"].sum()
```

```
101344.44467999999
```

```
df["life_exp"].sum() / df["life_exp"].count()
```

```
59.474439366197174
```

```
df["life_exp"].max()
```

```
82.603
```

```
# Sorting
```

```
df
```

	country	year	population	continent	life_exp	gdp_cap	
0	Afghanistan	1952	8425333	Asia	28.801	779.445314	
1	Afghanistan	1957	9240934	Asia	30.332	820.853030	
2	Afghanistan	1962	10267083	Asia	31.997	853.100710	
3	Afghanistan	1967	11537966	Asia	34.020	836.197138	
4	Afghanistan	1972	13079460	Asia	36.088	739.981106	
...	

```
df.sort_values("year", ascending=False).sort_values("life_exp") # wrong way
```

	country	year	population	continent	life_exp	gdp_cap	
1292	Rwanda	1992	7290203	Africa	23.599	737.068595	
0	Afghanistan	1952	8425333	Asia	28.801	779.445314	
552	Gambia	1952	284320	Africa	30.000	485.230659	
36	Angola	1952	4232095	Africa	30.015	3520.610273	
1344	Sierra Leone	1952	2143249	Africa	30.331	879.787736	
...	
1487	Switzerland	2007	7554661	Europe	81.701	37506.419070	
695	Iceland	2007	301931	Europe	81.757	36180.789190	
802	Japan	2002	127065841	Asia	82.000	28604.591900	
671	Hong Kong, China	2007	6980412	Asia	82.208	39724.978670	
803	Japan	2007	127467972	Asia	82.603	31656.068060	

1704 rows x 6 columns

```
df.sort_values(["year", "life_exp"], ascending=[False, True])
```


	country	year	population	continent	life_exp	gdp_cap
1463	Swaziland	2007	1133066	Africa	39.613	4513.480643
1043	Mozambique	2007	19951656	Africa	42.082	823.685621
1691	Zambia	2007	11746035	Africa	42.384	1271.211593



Creating the dataframe from scratch

```
pd.Series([10, 20, 30], index=[2, 3, 4])
```

```
2    10
3    20
4    30
dtype: int64
```

```
1000    Netherlands    1952    10301900    Europe    72.150    6341.571650
```

Row-oriented way

```
1140    Norway    1952    3327728    Europe    72.670    10005.121720
```

```
pd.DataFrame([[10, 20], [30, 40]], columns=["A", "B"])
```

	A	B
0	10	20
1	30	40



```
pd.DataFrame([[10, 20]], columns=["A", "B"])
```

	A	B
0	10	20



Column-oriented way

```
pd.DataFrame({"A": [10, 30], "B": [20, 40]})
```

	A	B
0	10	20
1	30	40



Optional Discussion

```
X = np.arange(12).reshape((3, 4))
row = np.array([0, 1, 2])
mask = np.array([1, 0, 1, 0], dtype=bool)
# print(X[row[:, np.newaxis], mask])
# didnt understand. how to interpret the slicing in print statement.
```

```
" Create mask, row 00 through 09 setting all game elements"
```

```
row
```

```
array([0, 1, 2])
```

```
X
```

```
array([[ 0,  1,  2,  3],
       [ 4,  5,  6,  7],
       [ 8,  9, 10, 11]])
```

```
row[:, np.newaxis]
```

```
array([[0],
       [1],
       [2]])
```

```
X[row[:, np.newaxis], mask]
```

```
array([[ 0,  2],
       [ 4,  6],
       [ 8, 10]])
```

```
row
```

```
array([0, 1, 2])
```

```
mask
```

```
array([ True, False,  True, False])
```

```
print(X[row[:, np.newaxis], mask])
```

```
[[ 0  2]
 [ 4  6]
 [ 8 10]]
```

```
print(X[row, mask])
```

```
-----
IndexError                                Traceback (most recent call last)
<ipython-input-208-75a392dca967> in <module>()
----> 1 print(X[row, mask])
```

```
IndexError: shape mismatch: indexing arrays could not be broadcast together
with shapes (3,) (2,)
```

SEARCH STACK OVERFLOW

<https://stackoverflow.com/questions/46124469/shape-mismatch-indexing-arrays-could-not-be->

```
row[:, np.newaxis]

array([[0],
       [1],
       [2]])

a = np.array((1,2,3, np.array([10,11])),dtype=object)
a

array([1, 2, 3, array([10, 11])], dtype=object)

b = a.copy()

np.shares_memory(a[3], b[3])

True

np.shares_memory(a[1], b[1])

False

np.shares_memory(a, b)

False

a[3][0] = 100

a

array([1, 2, 3, array([100, 11])], dtype=object)

b

array([1, 2, 3, array([100, 11])], dtype=object)
```

 0s completed at 00:07